

1. Online Speech-to-Text (ASR)

1. Online Speech-to-Text (ASR)

1. Concept Introduction

1.1 What is "ASR"?

1.2 Implementation Principles

1. Acoustic Model

2. Language Model

3. Pronunciation Lexicon

4. Decoder

5. End-to-End ASR

2. Project Architecture

Key Code

1. Speech Processing and Recognition Core (`targetmodel/targetmodel/asr.py`)

2. VAD Smart Recording (`targetmodel/targetmodel/asr.py`)

Code Analysis

3. Practical Operation

3.1 Configure Online ASR

3.2 Start and Test the Function

1. Concept Introduction

1.1 What is "ASR"?

ASR (Automatic Speech Recognition) is a technology that converts human speech signals into text. It is widely used in smart assistants, voice command control, automated customer service, real-time subtitle generation, and other fields. The goal of ASR is to enable machines to "understand" human language and convert it into a form that computers can process and understand.

1.2 Implementation Principles

The implementation of an ASR system mainly relies on the following key technical components:

1. Acoustic Model

- The acoustic model is responsible for converting the input sound signal into phonemes or sub-word units. This usually involves feature extraction steps, such as using Mel-frequency cepstral coefficients (MFCCs) or filter banks to represent the audio signal.
- These features are then fed into a deep neural network (DNN), convolutional neural network (CNN), recurrent neural network (RNN), or a more advanced Transformer architecture for training to learn how to map from audio features to corresponding phonemes or words.

2. Language Model

- The language model is used to predict the most likely next word given the preceding context, thereby helping to improve recognition accuracy. It is trained on a large amount of text data and understands which word sequences are more likely to occur.
- Common language models include n-gram models, RNN-based LMs, and the recently popular Transformer-based LMs.

3. Pronunciation Lexicon

- The pronunciation lexicon provides a mapping between words and their corresponding pronunciations. This is crucial for connecting the acoustic model and the language model, as it allows the system to understand and match the sounds it hears based on known pronunciation rules.

4. Decoder

- The decoder's task is to find the most likely sequence of words as output, given the acoustic model, language model, and pronunciation lexicon. This process usually involves complex search algorithms, such as the Viterbi algorithm or graph-based search methods, to find the optimal path.

5. End-to-End ASR

- With the development of deep learning, end-to-end ASR systems have emerged. These systems attempt to learn the text output directly from the raw audio signal without the need for separate designs of acoustic models, pronunciation lexicons, and language models. Such systems are often based on a sequence-to-sequence (Seq2Seq) framework, for example, using an attention mechanism or a Transformer architecture, which greatly simplifies the complexity of traditional ASR systems.

In general, modern ASR systems achieve efficient and accurate human speech-to-text capabilities by combining the above components and using large datasets and powerful computing resources for training. As technology advances, the performance of ASR systems continues to improve, and their application scenarios are becoming more widespread.

2. Project Architecture

Key Code

1. Speech Processing and Recognition Core (`largemodel/largemodel/asr.py`)

```
# From largemodel/largemodel/asr.py
def kws_handler(self)->None:
    if self.stop_event.is_set():
        return

    if self.listen_for_speech(self.mic_index):
        asr_text = self.ASR_conversion(self.user_speechdir) # Perform ASR
        conversion
        if asr_text == 'error': # Check if ASR result length is less than 4
            characters
            self.get_logger().warn("I still don't understand what you mean.
            Please try again")
            playsound(self.audio_dict[self.error_response]) # Error response
        else:
            self.get_logger().info(asr_text)
            self.get_logger().info("okay😊, let me think for a moment...")
            self.asr_pub_result(asr_text) # Publish ASR result
    else:
        return
```

```

def ASR_conversion(self, input_file:str)->str:
    if self.use_oline_asr:
        result=self.modelinterface.oline_asr(input_file)
        if result[0] == 'ok' and len(result[1]) > 4:
            return result[1]
        else:
            self.get_logger().error(f'ASR Error:{result[1]}') # ASR error.
            return 'error'
    else:
        result=self.modelinterface.SenseVoiceSmall_ASR(input_file)
        if result[0] == 'ok' and len(result[1]) > 4:
            return result[1]
        else:
            self.get_logger().error(f'ASR Error:{result[1]}') # ASR error.
            return 'error'

```

2. VAD Smart Recording (largemodel/largemodel/asr.py)

```

# From largemodel/largemodel/asr.py
def listen_for_speech(self,mic_index=0):
    p = pyaudio.PyAudio() # Create PyAudio instance.
    audio_buffer = [] # Store audio data.
    silence_counter = 0 # Silence counter.
    MAX_SILENCE_FRAMES = 90 # Stop after 900ms of silence (30 frames * 30ms)
    speaking = False # Flag indicating speech activity.
    frame_counter = 0 # Frame counter.
    stream_kwargs = {
        'format': pyaudio.paInt16,
        'channels': 1,
        'rate': self.sample_rate,
        'input': True,
        'frames_per_buffer': self.frame_bytes,
    }
    if mic_index != 0:
        stream_kwargs['input_device_index'] = mic_index

    # Prompt the user to speak via the buzzer.
    self.pub_beep.publish(UInt16(data = 1))
    time.sleep(0.5)
    self.pub_beep.publish(UInt16(data = 0))

    try:
        # Open audio stream.
        stream = p.open(**stream_kwargs)
        while True:
            if self.stop_event.is_set():
                return False

            frame = stream.read(self.frame_bytes, exception_on_overflow=False) #
            Read audio data.
            is_speech = self.vad.is_speech(frame, self.sample_rate) # VAD
            detection.

            if is_speech:

```

```

        # Detected speech activity.
        speaking = True
        audio_buffer.append(frame)
        silence_counter = 0
    else:
        if speaking:
            # Detect silence after speech activity.
            silence_counter += 1
            audio_buffer.append(frame) # Continue recording buffer.

            # End recording when silence duration meets the threshold.
            if silence_counter >= MAX_SILENCE_FRAMES:
                break
        frame_counter += 1
        if frame_counter % 2 == 0:
            self.get_logger().info('1' if is_speech else '-')
            # Real-time status display.
    finally:
        stream.stop_stream()
        stream.close()
        p.terminate()

    # Save valid recording (remove trailing silence).
    if speaking and len(audio_buffer) > 0:
        # Trim the last silent part.
        clean_buffer = audio_buffer[:-MAX_SILENCE_FRAMES] if len(audio_buffer) >
MAX_SILENCE_FRAMES else audio_buffer

        with wave.open(self.user_speechdir, 'wb') as wf:
            wf.setnchannels(1)
            wf.setsampwidth(p.get_sample_size(pyaudio.paInt16))
            wf.setframerate(self.sample_rate)
            wf.writeframes(b''.join(clean_buffer))
        return True

```

Code Analysis

The ASR (Speech-to-Text) function is provided by the `ASRNode` node (`asr.py`). This node is responsible for recording, converting, and publishing audio.

1. Audio Recording (`listen_for_speech`):

- This function uses the `pyaudio` library to capture the audio stream from the microphone.
- It integrates the `webrtcvad` library for voice activity detection (VAD). The function loops to read audio frames and uses `vad.is_speech()` to determine if each frame contains human voice.
- When speech is detected, the data is written to a buffer. When continuous silence is detected (defined by `MAX_SILENCE_FRAMES`), recording stops.
- Finally, the audio data in the buffer is written to a `.wav` file with the path `self.user_speechdir`.

2. Backend Selection and Execution (`ASR_conversion`):

- The `kws_handler` function calls the `ASR_conversion` function after a successful recording.
- This function determines which backend implementation to call by reading the ROS parameter `use_online_asr` (a boolean value).
- If it is `false`, the `self.modelinterface.SenseVoiceSmall_ASR` method is called, corresponding to local recognition.
- If it is `true`, the `self.modelinterface.online_asr` method is called, corresponding to online recognition.
- The function passes the audio file path as a parameter to the selected method and processes the returned result.

3. Result Publishing (`asr_pub_result`):

- After `ASR_conversion` returns valid text, `kws_handler` calls the `asr_pub_result` function.
- This function encapsulates the text string in a `std_msgs.msg.String` message and publishes it to the `/asr` topic through a ROS publisher.

3. Practical Operation

3.1 Configure Online ASR

1. Open the model interface configuration file `large_model_interface.yaml`:

```
vim ~/yahboom_ws/src/largemodel/config/large_model_interface.yaml
```

2. Fill in your API credentials:

Find the online ASR section, fill in your Key, and then select the speech recognition model.

```
# large_model_interface.yaml
# -----
--
# Large Language Models Settings
# -----
--
## Online Large Language Models
# Tongyi Platform Configuration
online_asr_api_key: "Your_AccessKey_ID"

# -----
--
# Speech Recognition Settings
# -----
--
## Online ASR
online_asr_model: 'paraformer-realtime-8k-v2' # Optional: 'paraformer-
realtime-v2', 'paraformer-realtime-v1', etc., using the API of the Tongyi
Qianwen platform
online_asr_sample_rate: 16000          # ASR audio sample rate
```

3. Open the main configuration file `yahboom.yaml`:

```
vim ~/yahboom_ws/src/largemodel/config/yahboom.yaml
```

4. Enable online ASR mode:

Set the `use_online_asr` parameter to `True`.

```
# yahboom.yaml

asr:
  ros__parameters:
    use_online_asr: True # Key: Change from False to True
    regional_setting : "China"
```

3.2 Start and Test the Function

1. Open a terminal and enter the start command:

```
ros2 launch largemodel asr_only.launch.py
```

2. Test:

Say to the microphone: "你好小亚". For international users, say: "Hi, Yahboom".

It will respond: "I'm here", and then you can start talking. Finally, the text converted from the recorded sound will be displayed in the terminal.

```
~/yahboom_ws$ ros2 launch largemodel asr_only.launch.py
[INFO] [launch]: All log files can be found below /home/sunrise/.ros/log/2025-08-07-15-40-45-012409-ubuntu-7665
[INFO] [launch]: Default logging verbosity is set to INFO
[INFO] [asr-1]: process started with pid [7672]
[asr-1] [INFO] [1754552481.072012346] [asr]: The asr model :SenseVoiceSmall is loaded
[asr-1] [INFO] [1754552481.517297362] [asr]: asr_node Initialization completed
[asr-1] [INFO] [1754552491.147085778] [asr]: I'm here
```